


ISTA 421 + INFO 521

Introduction to Machine Learning

Lecture 25:
Neural Networks II – The
Backpropagation Algorithm

Clay Morrison
 claytonm@email.arizona.edu
 Harvill 437A
 Phone 621-6609

22 November 2017

 1

Rest of the Semester

- Mon, Nov 20 – Intro to NN
- Wed, Nov 22 – NN 2: Backpropagation
- Mon, Nov 27 – NN 3: Sparse Autoencoders
- Wed, Nov 29 – Clustering: K-Means
- Mon, Dec 4 – Clustering: Gaussian Mixture Model
- Wed, Dec 6 – Principle Components Analysis

- Homework 5 (NN) due Thurs, Dec 7
- Final Project due Mon, Dec 11

- **Reminder: Please fill out course evaluation (TCE)**

Neural Networks and Deep Learning



Homework 5 and a Note about Notation

- **Tutorial:** Andrew Ng's *Unsupervised Feature Learning and Deep Learning* (UFLDL) tutorial on building a **sparse autoencoder** for classifying handwriting
 - Main focus: http://ufldl.stanford.edu/wiki/index.php/UFLDL_Tutorial
 - The first section on the Sparse Autoencoder is what we'll focus on.
 - There is also the newer UFLDL tutorial: <http://ufldl.stanford.edu/tutorial/>
- **Notation:**
 - The UFLDL tutorial (and in many cases in the neural network literature), matrices representing weights going from one layer to the next do so by having column indices for the *from* (source) and row indices for the *to* (destination) weight links.
 - Also, superscripts in parentheses are used to represent individuals in training data or layer index in the network; subscripts used to index into the elements of a vector/matrix.

The i^{th} input output value
 $(x^{(i)}, y^{(i)})$

The 3rd value of the i^{th} input vector
 $x_3^{(i)}$


scalar! (note the subscript)

$$\begin{array}{c} \text{From (source) layer activation} \\ \left\{ \begin{array}{c} \text{To (destination) layer} \\ \begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & & & \\ \hline & & & \\ \hline \end{array} \\ \hline \end{array} \right\} \cdot \left\{ \begin{array}{c} \text{From (source) layer activation} \\ \begin{array}{|c|} \hline x_2 \\ \hline \end{array} \end{array} \right\} = \left\{ \begin{array}{c} \text{To (destination) layer} \\ \begin{array}{|c|} \hline z_3^{(l+1)} \\ \hline \end{array} \end{array} \right\}
 \end{array}$$

$z^{(l+1)} = W^{(l)} x$
 vectors!
 matrix!



Feed Forward MLP Calculation (Forward Propagation)




n_l = number of layers = 3
 $W_{ij}^{(l)}$ = weight between unit j in layer l and unit i in layer $l + 1$
 $b_i^{(l)}$ = bias for unit i in layer $l + 1$
 $a_i^{(l)}$ = activation (output value) of unit i in layer l

$$\begin{aligned}
 a_1^{(2)} &= f(W_{11}^{(1)} x_1 + W_{12}^{(1)} x_2 + W_{13}^{(1)} x_3 + b_1^{(1)}) \\
 a_2^{(2)} &= f(W_{21}^{(1)} x_1 + W_{22}^{(1)} x_2 + W_{23}^{(1)} x_3 + b_2^{(1)}) \\
 a_3^{(2)} &= f(W_{31}^{(1)} x_1 + W_{32}^{(1)} x_2 + W_{33}^{(1)} x_3 + b_3^{(1)}) \\
 h_{W,b}(x) &= a_1^{(3)} = f(W_{11}^{(2)} a_1^{(2)} + W_{12}^{(2)} a_2^{(2)} + W_{13}^{(2)} a_3^{(2)} + b_1^{(2)})
 \end{aligned}$$

$$z_i^{(l)} = \sum_{j=1}^n W_{ij}^{(l-1)} x_j + b_i^{(l-1)} \text{ and } a_i^{(l)} = f(z_i^{(l)})$$



Feed Forward MLP Calculation (Forward Propagation)



n_l = number of layers = 3
 $W_{ij}^{(l)}$ = weight between unit j in layer l and unit i in layer $l + 1$
 $b_i^{(l)}$ = bias for unit i in layer $l + 1$
 $a_i^{(l)}$ = activation (output value) of unit i in layer l

Or, more compactly, if we make $f(\bullet)$ a vector function:

$$f([z_1, z_2, z_3]) = [f(z_1), f(z_2), f(z_3)]$$

$$z^{(2)} = W^{(1)} x + b^{(1)}$$


$$a^{(2)} = f(z^{(2)})$$

$$z^{(3)} = W^{(2)} a^{(2)} + b^{(2)}$$

$$h_{W,b}(x) = a^{(3)} = f(z^{(3)})$$



Feed Forward MLP Calculation (Forward Propagation)



Gradient Descent

- Given a loss function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} \left(\underbrace{\|h_{\theta}(x^{(i)}) - y^{(i)}\|}_{L_2 \text{ norm (i.e., Euclidean distance)}} \right)^2$$

when h_{θ} and y are vectors; when scalars, just subtract

- m is the number of examples and h some function of parameters θ
- Gradient descent updates the parameters in steps:

$$\theta_j = \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j}$$



Backpropagation Algorithm (for fully connected feedforward network)

training pairs: $\{(x^{(1)}, y^{(1)}), (x^{(m)}, y^{(m)})\}$

Define the loss

$$J(W, b) = \left[\frac{1}{m} \sum_{i=1}^m \left(\frac{1}{2} \|h_{W,b}(x^{(i)}) - y^{(i)}\|^2 \right) \right] + \underbrace{\frac{\lambda}{2} \sum_{l=1}^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (W_{ji}^{(l)})^2}_{\text{weight decay term}}$$

Implementation note: use `numpy.linalg.norm()`

Goal is to minimize $J(W, b)$

Neural networks can be prone to overfitting, so regularize

(note that reg term typically not over bias terms)

One iteration of gradient descent updates parameters W, b as follows:

$$W_{ij}^{(l)} = W_{ij}^{(l)} - \alpha \frac{\partial J(W, b)}{\partial W_{ij}^{(l)}} \quad \frac{\partial J(W, b)}{\partial W_{ij}^{(l)}} = \left[\frac{1}{m} \sum_{k=1}^m \frac{\partial J(W, b; x^{(k)}, y^{(k)})}{\partial W_{ij}^{(l)}} \right] + \lambda W_{ij}^{(l)}$$

$$b_i^{(l)} = b_i^{(l)} - \alpha \frac{\partial J(W, b)}{\partial b_i^{(l)}} \quad \frac{\partial J(W, b)}{\partial b_i^{(l)}} = \frac{1}{m} \sum_{k=1}^m \frac{\partial J(W, b; x^{(k)}, y^{(k)})}{\partial b_i^{(l)}}$$



9

Backpropagation Algorithm (for fully connected feedforward network)

Partial derivatives of cost function for single example (x, y)

Backprop Core (in words):


Given training example (x, y) ...

Run "forward pass" to compute all activations throughout network

For each node i in layer l , compute "error term" $\delta_i^{(l)}$ Measures how much node was "responsible" for any errors in output

For an output node, can directly measure the difference between network's activation and the true target value, to define $\delta_i^{(n_l)}$

For hidden units, compute $\delta_i^{(l)}$ based on weighted average of the error terms of the nodes that use $a_i^{(l)}$ as an input.



Backprop Core (in more detail):

(1) Perform feedforward pass, computing the activations for layers L_2, L_3 , and so on up to the output layer L_{n_l}

(2) For each output unit i in layer n_l (the output layer), set

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \frac{1}{2} \|y - h_{W,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)})$$

(3) For $l = n_l - 1, n_l - 2, n_l - 3, \dots, 2$

$$\text{For each node } i \text{ in layer } l, \text{ set } \delta_i^{(l)} = \left(\sum_{j=1}^{s_{l+1}} W_{ji}^{(l+1)} \delta_j^{(l+1)} \right) f'(z_i^{(l)})$$

Indices are correct!
i is layer l, j is layer l+1

(4) Compute the desired partial derivatives, which are given as:

$$\frac{\partial J(W, b; x, y)}{\partial W_{ij}^{(l)}} = a_j^{(l)} \delta_i^{(l+1)} \quad \frac{\partial J(W, b; x, y)}{\partial b_i^{(l)}} = \delta_i^{(l+1)}$$



10

Backpropagation Algorithm (for fully connected feedforward network)

Partial derivatives of cost function
for single example (x,y) $\frac{\partial J(W, b; x, y)}{\partial W_{ij}^{(l)}}$ $\frac{\partial J(W, b; x, y)}{\partial b_i^{(l)}}$

Backprop Core (Vectorized):

(1) Perform feedforward pass

(2) For the output layer (layer n_l), set $\delta^{(n_l)} = -(y - a^{(n_l)}) \bullet f'(z^{(n_l)})$

(3) For $l = n_l - 1, n_l - 2, n_l - 3, \dots, 2$, set $\delta^{(l)} = \left((W^{(l+1)})^\top \delta^{(l+1)} \right) \bullet f'(z^{(l)})$

(4) Compute the desired partial derivatives, which are given as:

$$\nabla_{W^{(l)}} J(W, b; x, y) = \delta^{(l+1)} (a^{(l)})^\top$$


Yes, this is an outer product!!
Implementation note: look at `numpy.linalg.outer()`

$$\nabla_{b^{(l)}} J(W, b; x, y) = \delta^{(l+1)}$$

If sigmoid activation:

$$f'(z_i^{(l)}) = a_i^{(l)}(1 - a_i^{(l)})$$

store during forward pass



Backprop Core (in more detail):

(1) Perform feedforward pass, computing the activations for layers L_2, L_3 , and so on up to the output layer L_{n_l}

(2) For each output unit i in layer n_l (the output layer), set

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \frac{1}{2} \|y - h_{W,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)})$$

(3) For $l = n_l - 1, n_l - 2, n_l - 3, \dots, 2$
For each node i in layer l , set $\delta_i^{(l)} = \left(\sum_{j=1}^{s_{l+1}} W_{ji}^{(l+1)} \delta_j^{(l+1)} \right) f'(z_i^{(l)})$

(4) Compute the desired partial derivatives, which are given as:

$$\frac{\partial J(W, b; x, y)}{\partial W_{ij}^{(l)}} = a_j^{(l)} \delta_i^{(l+1)} \quad \frac{\partial J(W, b; x, y)}{\partial b_i^{(l)}} = \delta_i^{(l+1)}$$



11

Backpropagation Algorithm (for fully connected feedforward network)

Partial derivatives of cost function
for single example (x,y) $\frac{\partial J(W, b; x, y)}{\partial W_{ij}^{(l)}}$ $\frac{\partial J(W, b; x, y)}{\partial b_i^{(l)}}$

Backprop Core (Vectorized):

(1) Perform feedforward pass

(2) For the output layer (layer n_l), set $\delta^{(n_l)} = -(y - a^{(n_l)}) \bullet f'(z^{(n_l)})$

(3) For $l = n_l - 1, n_l - 2, n_l - 3, \dots, 2$, set $\delta^{(l)} = \left((W^{(l+1)})^\top \delta^{(l+1)} \right) \bullet f'(z^{(l)})$

(4) Compute the desired partial derivatives, which are given as:


$$\nabla_{W^{(l)}} J(W, b; x, y) = \delta^{(l+1)} (a^{(l)})^\top$$

$$\nabla_{b^{(l)}} J(W, b; x, y) = \delta^{(l+1)}$$

If sigmoid activation:

$$f'(z_i^{(l)}) = a_i^{(l)}(1 - a_i^{(l)})$$

store during forward pass



The Complete Backprop Algorithm as Gradient Descent:

(1) Set $\Delta W^{(l)} := 0$, $\Delta b^{(l)} := 0$ for all l (note that ' ΔW ' and ' Δb ' are single variables)

(2) For $i = 1$ to m ,

(a) Use backpropagation (Backprop Core) to compute $\nabla_{W^{(l)}} J(W, b; x, y)$

(b) Set $\Delta W^{(l)} := \Delta W^{(l)} + \nabla_{W^{(l)}} J(W, b; x, y)$

(c) Set $\Delta b^{(l)} := \Delta b^{(l)} + \nabla_{b^{(l)}} J(W, b; x, y)$

(3) Update the parameters:

$$W^{(l)} = W^{(l)} - \alpha \left[\left(\frac{1}{m} \Delta W^{(l)} \right) + \lambda W^{(l)} \right] \quad b^{(l)} = b^{(l)} - \alpha \left[\frac{1}{m} \Delta b^{(l)} \right]$$

gradient gradient

Repeat (2) and (3) until "convergence"

Initializing Parameters

- If all parameters start off at identical values (e.g., 0), then all hidden layer units would learn the same function of the input.
- Random initialization breaks symmetry.
- Often sample small random values near zero:

$$\mathcal{N}(0, \epsilon^2) \quad (\text{e.g., } \epsilon = 0.01)$$

- Another heuristic:

$$\text{Uniform} \left[-\sqrt{\frac{6}{n_{\text{in}} + n_{\text{out}} + 1}}, \sqrt{\frac{6}{n_{\text{in}} + n_{\text{out}} + 1}} \right]$$



Improvements to Gradient Descent

$$W^{(l)} = W^{(l)} - \alpha \left[\left(\frac{1}{m} \Delta W^{(l)} \right) + \lambda W^{(l)} \right] \quad b^{(l)} = b^{(l)} - \alpha \left[\frac{1}{m} \Delta b^{(l)} \right]$$

- Automatically adapting α
- Newton's method(s): using Hessian information
- L-BFGS
- Conjugate Gradient

What you need to provide:

A function that returns the cost ($J(\theta)$) and grad ($\nabla_{\theta} J(\theta)$)

Initial theta

```
J = lambda x: your_fn(x, vis_size, hid_size, lambda_, data)
scipy.optimize.minimize(fun=J, x0=theta, method='L-BFGS-B')
```

Debugging: Gradient Checking

Numerically checking derivatives

Recall mathematical definition of the derivative (assuming J is fn of θ):

$$\frac{\partial J(\theta)}{\partial \theta} = \lim_{\epsilon \rightarrow 0} \frac{J(\theta + \epsilon) - J(\theta - \epsilon)}{2\epsilon}$$

At any specific value of θ we can numerically approximate the derivative by:

$$\frac{J(\theta + \text{EPSILON}) - J(\theta - \text{EPSILON})}{2 \times \text{EPSILON}}$$

In practice, set EPSILON to a small constant, such as $10^{-4} = 0.0001$.

You can use this numerical approximation to compare against a function $g(\theta) = \frac{\partial J(\theta)}{\partial \theta}$ used to compute the gradient.

$$g(\theta) \approx \frac{J(\theta + \text{EPSILON}) - J(\theta - \text{EPSILON})}{2 \times \text{EPSILON}}$$

With EPSILON = 10^{-4} , you'll usually get a similarity between $g(\theta)$ and your numerical approximation to at least 4 significant digits.



Debugging: Gradient Checking

Numerically checking derivatives

$$g(\theta) \approx \frac{J(\theta + \text{EPSILON}) - J(\theta - \text{EPSILON})}{2 \times \text{EPSILON}}$$

$$g(\theta) = \frac{\partial J(\theta)}{\partial \theta}$$

The above approximates the gradient for a single parameter θ .

If we are computing the derivative of each component parameter i in vector θ , then we can modify the above estimate by varying the specific parameter value by EPSILON:

$$g_i(\theta) = \frac{\partial J(\theta)}{\partial \theta_i} \quad \theta = \begin{bmatrix} \theta_1 \\ \theta_2 \\ \vdots \\ \theta_i \\ \vdots \\ \theta_N \end{bmatrix} \quad \mathbf{e}_i = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \\ \vdots \\ 0 \end{bmatrix}$$

Let $\theta^{(i+)} = \theta + \text{EPSILON} \times \mathbf{e}_i$

$\theta^{(i-)} = \theta - \text{EPSILON} \times \mathbf{e}_i$

The comparison then becomes:

$$g_i(\theta) \approx \frac{J(\theta^{(i+)}) - J(\theta^{(i-)})}{2 \times \text{EPSILON}}$$

The goal of each iteration of the Backpropagation Algorithm is to compute the gradient

$$W^{(l)} = W^{(l)} - \alpha \left[\underbrace{\left(\frac{1}{m} \Delta W^{(l)} \right)}_{\text{gradient}} + \lambda W^{(l)} \right] \quad b^{(l)} = b^{(l)} - \alpha \left[\underbrace{\left(\frac{1}{m} \Delta b^{(l)} \right)}_{\text{gradient}} \right]$$

$$\nabla_{W^{(l)}} J(W, b) \quad \nabla_{b^{(l)}} J(W, b)$$

The above numerical gradient estimation can be used to numerically compute the derivatives of $J(W, b)$ and compare to the computations of $\nabla_{W^{(l)}} J(W, b)$, $\nabla_{b^{(l)}} J(W, b)$